

References

- [1] “Amax high performance computing white papers,” <http://www.amax.com/hpc/whitepapers.asp>.
- [2] “Cusp-library – generic parallel algorithms for sparse matrix and graph computations,” <http://code.google.com/p/cusp-library>.
- [3] “<http://tbp.berkeley.edu/jdonald/research/cmp/>.”
- [4] “Intel’s 48-core Single-Chip Cloud Computer,” <http://www.intel.com/content/www/us/en/research/intel-labs-single-chip-cloud-computer.html>.
- [5] “Nvidia research programs and centers,” <http://research.nvidia.com/>.
- [6] “NVIDIA Tesla’s GPU computing solutions,” http://www.nvidia.com/object/tesla_computing_solutions.html.
- [7] “NVIDIA Tesla’s Servers and Workstations,” <http://www.nvidia.com/object/tesla-servers.html>.
- [8] “Openacc directives for accelerators,” <http://openacc-standard.org>.
- [9] “SimpleScalar tool,” <http://www.simplescalar.com>.
- [10] “Tilera’s 64-core Processor: TILEProf64,” http://www.tilera.com/products/processors/TILE-Gx_Family.
- [11] “Top 500 supercomputers,” <http://www.top500.org/>.
- [12] “NVIDIA CUBLAS library reference v2.3,” 2009.
- [13] “NVIDIA CUDA architecture v1.1,” 2009.
- [14] “NVIDIA CUDA Programming Guide v2.3,” 2009.
- [15] “NVIDIA CUDA reference manual v2.3,” 2009.
- [16] “NVIDIA’s next generation CUDA compute architecture: Fermi,” 2009, white paper.
- [17] “NVIDIA CUDA Programming Guide v3.2,” Dec 2010, http://developer.download.nvidia.com/compute/cuda/3.2_prod/toolkit/docs/CUBLAS_Library.pdf.

- [18] A. Agarwal, D. Blaauw, V. Zolotov, S. Sundareswaran, M. Zhao, K. Gala, and R. Panda, “Path-based statistical timing analysis considering inter- and intra-die correlations,” in *Proceedings of the TAU Conference*, 2002.
- [19] A. Agarwal, B. C. Paul, and K. Roy, “A novel fault tolerant cache to improve yield in nanometer technologies,” in *IEEE International On-Line Testing Symposium*, 2004, pp. 149–154.
- [20] G. Alaghband, “Parallel sparse matrix solution and performance,” *Parallel Comput.*, vol. 21, no. 9, pp. 1407–1430, 1995.
- [21] AMD Inc., “Multi-core processors—the next evolution in computing (White Paper),” 2006, <http://multicore.amd.com>.
- [22] —, “AMD developer center,” <http://developer.amd.com/GPU>, 2011.
- [23] —, “AMD Steam SDK,” <http://developer.amd.com/gpu/ATISStreamSDK>, 2011.
- [24] —, “AMD Steam SDK,” <http://developer.amd.com/gpu/ATISStreamSDK>, 2011.
- [25] M. M. Baskaran and R. Bordawekar, “Optimizing sparse matrix-vector multiplication on GPUs,” IBM Research Division, IBM Research Report RC24704, Apr. 2009.
- [26] N. Bell and M. Garland, “Efficient Sparse Matrix-Vector Multiplication on CUDA,” http://www.nvidia.com/object/nvidia_research_pub_001.html.
- [27] —, “Efficient sparse matrix-vector multiplication on CUDA,” NVIDIA Corporation, NVIDIA Technical Report NVR-2008-004, Dec. 2008.
- [28] —, “Implementing sparse matrix-vector multiplication on throughput-oriented processors,” in *Proceedings of the Conference on High Performance Computing Networking, Storage and Analysis*, ser. SC ’09. New York, NY, USA: ACM, 2009, pp. 18:1–18:11.

- [29] A. Bhavanagarwala, X. Tang, and J. Meindl, "The impact of intrinsic device fluctuations on cmos sram cell stability," in *IEEE Journal of Solid-State Circuits*, 2001, pp. 658–665.
- [30] S. Borkar, "Thousand core chips: a technology perspective," in *Proc. Design Automation Conf. (DAC)*, 2007, pp. 746–749.
- [31] S. Borkar, P. Dubey, K. Kahn, D. Kuck, H. Mulder, S. Pawlowski, and J. Rattner, "Platform 2015: Intel processor and platform evolution for the next decade," Intel white paper (http://tbp.berkeley.edu/jdonald/research/cmp/borkar_2015.pdf).
- [32] S. Borkar, T. Karnik, and V. De, "Design and reliability challenges in nanometer technologies," in *Proc. Design Automation Conf. (DAC)*. IEEE Press, 2004, pp. 75–75.
- [33] S. Borkar, T. Karnik, S. Narendra, J. Tschanz, A. Keshavarzi, and V. De, "Parameter variations and impact on circuits and microarchitecture," in *Proc. Design Automation Conf. (DAC)*. IEEE Press, 2003, pp. 338–342.
- [34] K. Bowman, S. G. Duvall, and J. D. Meindl, "Impact of die-to-die and within-die parameter fluctuations on the maximum clock frequency distribution," in *International Solid-State Circuits Conference*, Feb 2001, pp. 278–279.
- [35] L. Buatois, G. Caumon, and B. Levy, "Concurrent number cruncher: a GPU implementation of a general sparse linear solver," *Int. J. Parallel Emerg. Distrib. Syst.*, vol. 24, no. 3, pp. 205–223, 2009.
- [36] T. C. Chen, "Where CMOS is going: trendy hype vs. real technology," in *International Solid-State Circuits Conference*, Feb. 2006, pp. 22–27.
- [37] J. Clabes and et al., "Design and implementation of the power5 microprocessor," Feb. 2004, pp. 56–57.
- [38] R. Couturier, "Sparse systems solving on GPUs with GMRES," *The Journal of Supercomputing*, 2011.
- [39] R. Couturier and S. Domas, "Sparse systems solving on GPUs with GMRES," *The Journal of Supercomputing*, pp. 1–13–13, 2011. [Online]. Available: <http://www.springerlink.com/content/w426327851060822/>

- [40] Y. Deng, B. Wang, and S. Mu, “Taming irregular EDA applications on GPUs,” in *Computer-Aided Design - Digest of Technical Papers, 2009. ICCAD 2009. IEEE/ACM International Conference on*, 2009, pp. 539–546.
- [41] S. G. Duvall, “Statistical circuit modeling and optimization,” in *Intl. Workshop Statistical Metrology*, Jun 2000, pp. 56–63.
- [42] P. Friedberg, Y. Cao, J. Cain, R. Wang, J. Rabaey, and C. Spanos, “Modeling within-die spatial correlation effects for process design co-optimization,” in *Proceedings of the 6th International Symposium on Quality of Electronic Design*, 2005, pp. 516–521.
- [43] D. G  ddeke, “General-purpose computation using graphics hardware,” <http://www.gpgpu.org/>, 2011.
- [44] K. Gulati and S. P. Khatri, *Hardware Acceleration of EDA Algorithms*. Springer, 2010.
- [45] R. J. Hanson and D. R. Kincaid, “Notes on GMRES algorithm organization,” Dept. of Computer Sciences, University of Texas at Austin, Technical Report TR-05-05, Mar. 2005.
- [46] L. He, W. Liao, and M. R. Stan, “System level leakage reduction considering the interdependence of temperature and leakage,” in *Proc. Design Automation Conf. (DAC)*, 2004, pp. 12–17.
- [47] M. Horiguchi, “Redundancy techniques for high-density drams,” in *Second Annual IEEE International Conference on Innovative Systems in Silicon*, Oct 1997, pp. 22–29.
- [48] L. Hsu, R. Iyer, S. Makineni, S. Reinhardt, and D. Newell, “Exploring the cache design space for large-scale CMPs,” *ACM SIGARCH Computer Architecture News*, vol. 33, no. 4, pp. 24–33, Nov. 2005.
- [49] S. K. Hsu, N. S. Kim, and S.-L. Lu, “Memory array circuits for computer architects,” in *HPCA Tutorial*, 2006.
- [50] Intel Corporation, “Intel multi-core processors, making the move to quad-core and beyond (White Paper),” 2006, <http://www.intel.com/multi-core>.
- [51] —, “Intel Core 2 Extreme Process QX 9000 Series,” 2009, <http://download.intel.com/design/processor/datashts/318726.pdf>.

- [52] J. K. John, J. S. Hu, and S. G. Ziavras, “Optimizing the thermal behavior of subarrayed data caches,” in *Proc. IEEE Int. Conf. on Computer Design (ICCD)*, 2005.
- [53] K. Kang, B. C. Paul, and K. Roy, “Statistical timing analysis using levelized covariance propagation considering systematic and random variations of process parameters,” in *ACM Trans. on Design Automation of Electronics Systems*, vol. 11, Oct 2006, pp. 848–879.
- [54] N. Kapre and A. DeHon, “Performance comparison of single-precision SPICE model-evaluation on FPGA, GPU, cell, and multi-core processors,” in *International Conference on Field Programmable Logic and Applications*, Sept. 2009, pp. 65–72.
- [55] Khronos Group, “Open Computing Language (OpenCL),” <http://www.khronos.org/opencl>, 2011.
- [56] D. B. Kirk and W.-M. Hwu, *Programming Massively Parallel Processors: A Hands-on Approach*. San Francisco, CA: Morgan Kaufmann Publishers Inc., 2010.
- [57] —, *Programming Massively Parallel Processors: A Hands-on Approach, 2ed.* San Francisco, CA: Morgan Kaufmann Publishers Inc., 2013.
- [58] J. U. Knickerbocker, P. S. Andry, B. Dang, R. R. Horton, C. S. Patel, R. J. Polastre, K. Sakuma, E. S. Sprogis, C. K. Tsang, B. C. Webb, and S. L. Wright, “3D Silicon Integration,” in *2008 58th Electronic Components and Technology Conference (ECTC 2008)*. IEEE, 2008, pp. 538–543.
- [59] P. Kongetira, K. Aingaran, and K. Olukotun, “Niagara: A 32-Way multithreaded sparcs processor,” *IEEE Micro*, vol. 25, no. 2, pp. 21–29, 2005.
- [60] J. C. Ku, S. Ozdemir, G. Memik, and Y. Ismail, “Thermal management of on-chip caches through power density minimization,” in *International Symposium on Microarchitecture (MICRO)*, 2005.
- [61] R. Kumar, V. Zyuban, and D. Tullsen, “Interconnections in multi-core architectures: Understanding mechanisms, overheads and scaling,” in *International Symposium on Computer Architecture*, June 2005, pp. 408–419.

- [62] R. Li and Y. Saad, “GPU-accelerated preconditioned iterative linear solvers,” *The Journal of Supercomputing*, vol. 63, no. 2, pp. 443–466, 2010. [Online]. Available: <http://static.msi.umn.edu/rreports/2010/112.pdf>
- [63] Y. Li, D. Brooks, Z. Hu, and K. Skadron, “Performance, energy, and thermal considerations for SMT and CMP architectures,” in *Proc. IEEE Int. Symp. on High Performance Computer Architecture (HPCA)*, 2005, pp. 71–82.
- [64] Y. Li, B. C. Lee, D. Brooks, Z. Hu, and K. Skadron, “CMP design space exploration subject to physical constraints,” *Proc. IEEE Int. Symp. on High Performance Computer Architecture (HPCA)*, pp. 15–26, 2006.
- [65] P. Liu, Z. Qi, H. Li, L. Jin, W. Wu, S. X.-D. Tan, and J. Yang, “Fast thermal simulation for architecture level dynamic thermal management,” in *Proc. Int. Conf. on Computer Aided Design (ICCAD)*, Nov 2005, pp. 639–644.
- [66] W. Liu, X. Jin, J. Chen, M.-C. Jeng, Z. Liu, Y. Cheng, K. Chen, M. Chan, K. Hui, J. Huang, R. Tu, P. Ko, and C. Hu, “Bsim3v3.2.1 mosfet model users’ manual,” EECS Department, University of California, Berkeley, Tech. Rep. UCB/ERL M99/19, 1999. [Online]. Available: <http://www.eecs.berkeley.edu/Pubs/TechRpts/1999/3629.html>
- [67] C. McNairy and R. Bhatia, “Montecito: A dual-Core, dual-Thread itanium processor,” *IEEE Micro*, vol. 25, no. 2, pp. 10–20, 2005.
- [68] K. Meng and R. Joseph, “Process variation aware cache leakage management,” in *Proceedings of the International Symposium on Low Power Electronics and Design*, 2006, pp. 262–267.
- [69] M. Monchiero, R. Canal, and A. Gonzalez, “Power/Performance/Thermal design-Space exploration for multicore architectures,” *IEEE Trans. On Parallel and Distributed Systems*, vol. 19, no. 5, pp. 666–680, 2008.
- [70] T. Mudge, “Power: A first class design constraint for future architectures,” in *Proc. IEEE Int. Symp. on High Performance Computer Architecture (HPCA)*, vol. 1970, Dec. 2000, pp. 215–224.

- [71] S. Mukhopadhyay, H. Mahmoodi, and K. Roy, “Modeling and estimation of failure probability due to parameter variations in nano-scale srams for yield enhancement,” in *Symp. on VLSI Circuits*, 2004.
- [72] M. Naumov, “Parallel solution of sparse triangular linear systems in the preconditioned iterative methods on the GPU,” NVIDIA Corp., NVIDIA Technical Report NVR-2011-001, June 2011.
- [73] E. G. Ng and B. W. Peyton, “Block sparse cholesky algorithms on advanced uniprocessor computers,” *SIAM J. Sci. Comput.*, vol. 14, no. 5, pp. 1034–1056, 1993.
- [74] R. R. P. V. Nooyen, C. Vuik, C. Vuik, and P. Wesseling, “Some parallelizable preconditioners for gmres,” 1996.
- [75] NVIDIA, “CUBLAS library,” February 2011, <http://developer.nvidia.com/cuBLAS>.
- [76] —, “CUSPARSE library,” January 2011, <http://developer.nvidia.com/cuSPARSE>.
- [77] —, “CUDA C programming guide,” docs.nvidia.com/cuda/pdf/CUDA_C_Programming_Guide.pdf, October 2012.
- [78] NVIDIA Corporation, “CUBLAS library v5.0,” <https://developer.nvidia.com/cublas>.
- [79] NVIDIA Corporation., “Tesla C2050 / C2070 GPU computing processor,” http://www.nvidia.co.uk/object/product_tesla_C2050_C2070_uk.html.
- [80] NVIDIA Corporation, 2011, <http://www.nvidia.com>.
- [81] NVIDIA Corporation, “CUDA (Compute Unified Device Architecture),” 2011, http://www.nvidia.com/object/cuda_home.html.
- [82] NVIDIA Corporation, “CUSPARSE library v5.0,” October 2012, <http://developer.nvidia.com/cuSPARSE>.
- [83] M. Okuno and et al, “45-nm node cmos integration with a novel sti structure and full-ncs/cu interlayers for low-operation-power (lop) applications,” in *IEEE International Electron Devices Meeting*, Dec 2005, pp. 52–55.

- [84] “OpenMP,” <http://www.openmp.org/blog/>.
- [85] M. Orshansky, L. Milor, and C. Hu, “Characterization of spatial intrafield gate cd variability, its impact on circuit performance, and spatial mask-level correction,” in *IEEE Trans. on Semiconductor Devices*, vol. 17, no. 1, Feb 2004, pp. 2–11.
- [86] D. Payne, “GPU versus multicore in EDA,” <http://www.chipdesignmag.com/>.
- [87] J. Reinders, *Intel Threading Building Blocks: Outfitting C++ for Multi-core Processor Parallelism (Paperback)*. O’Reilly Media, 2007, ISBN 978-0-596-51480-8.
- [88] E. Rothberg and A. Gupta, “Techniques for improving the performance of sparse matrix factorization on multiprocessor workstations,” in *Supercomputing ’90: Proceedings of the 1990 ACM/IEEE conference on Supercomputing*. Los Alamitos, CA, USA: IEEE Computer Society Press, 1990, pp. 232–241.
- [89] K. Sakuma, P. S. Andry, C. K. Tsang, S. L. Wright, B. Dang, C. S. Patel, B. C. Webb, J. Maria, E. J. Sprogis, S. K. Kang, R. J. Polastre, R. R. Horton, and J. U. Knickerbocker, “3D chip-stacking technology with through-silicon vias and low-volume lead-free interconnections,” *IBM Journal of Research and Development*, vol. 52, no. 6, pp. 611–622, 2008.
- [90] J. Sanders and E. Kandrot, *CUDA by Example: An Introduction to General-Purpose GPU Programming*. Addison Wesley, 2010.
- [91] K. Sankaranarayanan, S. Velusamy, M. Stan, and K. Skadron, “A case for thermal-aware floorplanning at the microarchitectural level,” in *The Journal of Instruction-Level Parallelism*, Oct. 2005, pp. 1–16.
- [92] R. S. Shelar and S. S. Sapatnekar, “An efficient algorithm for low power pass transistor logic synthesis,” in *VLSI Design*, 2002, pp. 87–92.
- [93] T. Sherwood, E. Perelman, and B. Calder, “Basic block distribution analysis to find periodic behavior and simulation points in applications,” in *International Conference on Parallel Architecture and Compilation Techniques*, 2001.
- [94] P. Shivakumar and N. P. Jouppi, “CACTI 3.0: An integrated cache timing, power, and area model,” 2001.

- [95] B. Sinharoy, R. Kalla, J. Tendler, R. Eickemeyer, and J. Joyner, "Power5 system microarchitecture," *IBM J. Research and Development*, vol. 49, no. 4, pp. 505–521, 2005.
- [96] J. Srinivasan, S. V. Adve, P. Bose, and J. A. Rivers, "The case for microarchitectural awareness of lifetime reliability," in *International Symposium on Computer Architecture*, 2004.
- [97] —, "The impact of technology scaling on lifetime reliability," in *Dependable Systems and Networks*, 2004.
- [98] —, "Exploiting structural duplication for lifetime reliability enhancement," in *International Symposium on Computer Architecture*, 2005, pp. 520–531.
- [99] V. Szekely, "Identification of RC networks by deconvolution: Chances and limits," *IEEE Trans. on Circuits and Systems I: Fundamental Theory and Applications*, vol. 45, no. 3, pp. 244–258, 1998.
- [100] R. Teodorescu, B. Greskamp, J. Nakano, S. R. Sarangi, A. Tiwari, and J. Torrellas, "A model of parameter variation and resulting timing errors for microarchitects," in *Workshop on Architectural Support for Gigascale Integration (ASGI)*, Jun 2007.
- [101] S. Tomov *et al.*, "Dense Linear Algebra Solvers for Multicore with GPU Accelerators," in *Distributed Processing, Workshops and Phd Forum (IPDPSW)*, 2010, pp. 1–8.
- [102] S. Tomov, R. Nath, H. Ltaief, and J. Dongarra, "Dense Linear Algebra Solvers for Multicore with GPU Accelerators," in *Distributed Processing, Workshops and Phd Forum (IPDPSW)*. IEEE, 2010, pp. 1–8.
- [103] Y.-F. Tsai, Y.-F. Tsai, A. H. Ankadi, V. Narayanan, and M. J. Irwin, "Chippower: an architecture-level leakage simulator," in *Proceedings of IEEE International SOC Conference*, 2004, pp. 395–398.
- [104] M. Wang, H. Klie, *et al.*, "Solving sparse linear systems on NVIDIA Tesla GPUs," in *Proc of the 9th Intl. Conf. on Computational Science*, 2009, pp. 864–873.
- [105] P. Wijetunga, "High-performance crossbar design for system-on-chip," in *International Workshop on System-on-Chip for Real-Time Applications*, 2003, pp. 138–143.

- [106] W. Wu, J. Yang, S. X.-D. Tan, and S.-L. Lu, “Improving the reliability of on-chip caches under process variations,” in *Proc. IEEE Int. Conf. on Computer Design (ICCD)*, 2007, pp. 325–332.
- [107] S. Xu, W. Xue, K. Wang, and H. Lin, “Generating approximate inverse preconditioners for sparse matrices using CUDA and GPGPU,” *Journal of Algorithm & Computational Technology*, vol. 5, no. 3, pp. 475–500, 2010.